

##BITS F464 - Semester 1 - MACHINE LEARNING

ASSIGNMENT 1 - LINEAR MODELS FOR REGRESSION AND CLASSIFICATION

Group 9

Tarun Raman, Jayesh Totlani, Ananya Jain, Anamay Trivedi, Divyanshu

2021AAPS2308H, 2021AAPS1815H, 2021AAPS2121H 2021AAPS2217H, 2021AAPS1937H

This assignment aims to identify the differences between three sets of Machine Learning models.

1. Dataset Generation

Imported the included dataset 'diabetes2.csv' which has been attached along in the submission. Synthetic Data Synthesizer code has been provided in a commented way.

```
#!/usr/bin/env python
#%pip install sdv
#%pip install urllib3==1.26.7

#import sdv

#from sdv.datasets.local import load_csvs

# assume that my_folder contains 1 CSV file named 'guests.csv'
#datasets = load_csvs(folder_name='/content')

# the data is available under the file name
#guests_table = datasets['diabetes2']

#from sdv.metadata import SingleTableMetadata

#metadata = SingleTableMetadata()

#metadata.detect_from_csv(filepath='/content/diabetes2.csv')

#from sdv.lite import SingleTablePreset

#synthesizer = SingleTablePreset(
#    metadata,
#    name='FAST_ML'
```

```

#)

#synthesizer.fit(
#    data=guests_table

#)

#synthetic_data = synthesizer.sample(
#    num_rows=500
#)

#synthetic_data.head()

#synthetic_data.to_csv('diabetes_new.csv')

# This is formatted as code

```

2. Preprocess and perform exploratory data analysis of the dataset obtained

Importing Libraries

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import random

```

Data Analysis & Pre-Processing

```

df=pd.read_csv(r"diabetes_new.csv")
df=df.drop("Unnamed: 0",axis=1)
df.head()

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
BMI \					
0	6	139	75	35	159
35.448240					
1	4	79	42	43	177
28.084784					
2	1	161	79	29	202
38.425897					
3	5	147	62	22	122
26.103594					
4	1	136	83	45	169
40.988707					

	DiabetesPedigreeFunction	Age	Outcome
0	0.780996	37	1
1	0.335059	29	1
2	0.618149	35	1
3	0.147669	37	0
4	0.078000	21	0

df.corr()

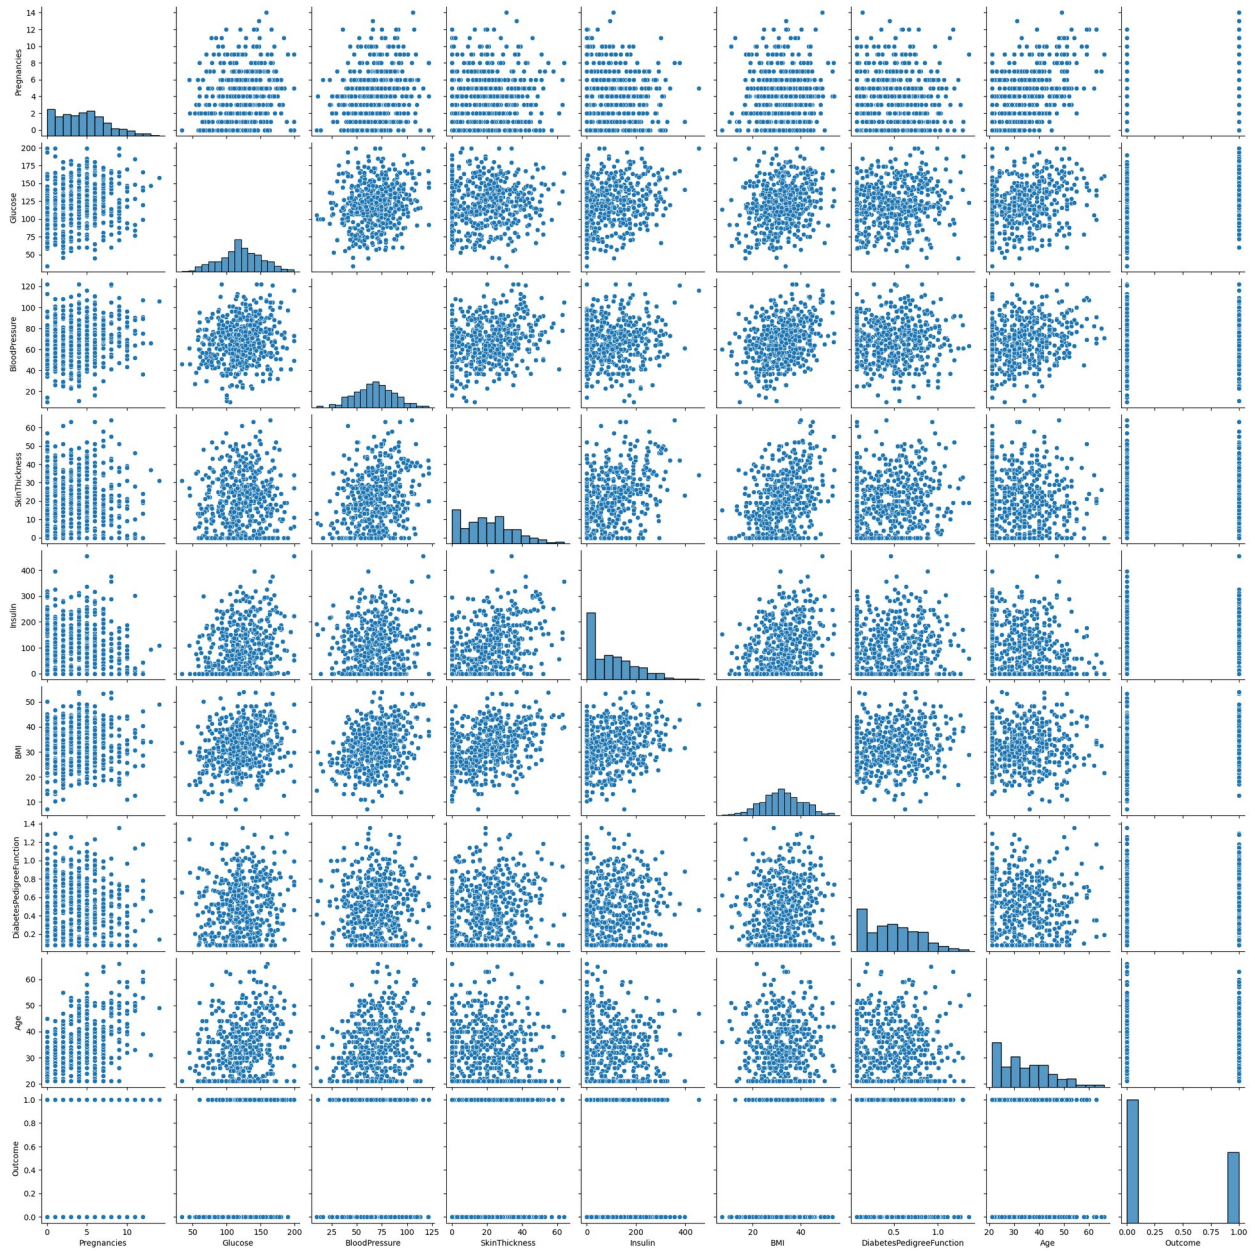
	Pregnancies	Glucose	BloodPressure	
SkinThickness \				
Pregnancies	1.000000	0.185623	0.174361	-
0.064803				
Glucose	0.185623	1.000000	0.196215	
0.056860				
BloodPressure	0.174361	0.196215	1.000000	
0.257591				
SkinThickness	-0.064803	0.056860	0.257591	
1.000000				
Insulin	-0.129844	0.252544	0.096254	
0.447154				
BMI	0.060411	0.233849	0.350713	
0.378895				
DiabetesPedigreeFunction	-0.056273	0.112083	0.029064	
0.135503				
Age	0.523230	0.246321	0.247437	-
0.087813				
Outcome	0.176085	0.351813	0.072567	
0.050403				

	Insulin	BMI	DiabetesPedigreeFunction
\			
Pregnancies	-0.129844	0.060411	-0.056273
Glucose	0.252544	0.233849	0.112083
BloodPressure	0.096254	0.350713	0.029064
SkinThickness	0.447154	0.378895	0.135503
Insulin	1.000000	0.261752	0.110587
BMI	0.261752	1.000000	0.152005
DiabetesPedigreeFunction	0.110587	0.152005	1.000000
Age	-0.134450	0.010781	-0.014336
Outcome	0.102682	0.271032	0.108521

	Age	Outcome
Pregnancies	0.523230	0.176085
Glucose	0.246321	0.351813
BloodPressure	0.247437	0.072567
SkinThickness	-0.087813	0.050403
Insulin	-0.134450	0.102682
BMI	0.010781	0.271032
DiabetesPedigreeFunction	-0.014336	0.108521
Age	1.000000	0.134397
Outcome	0.134397	1.000000

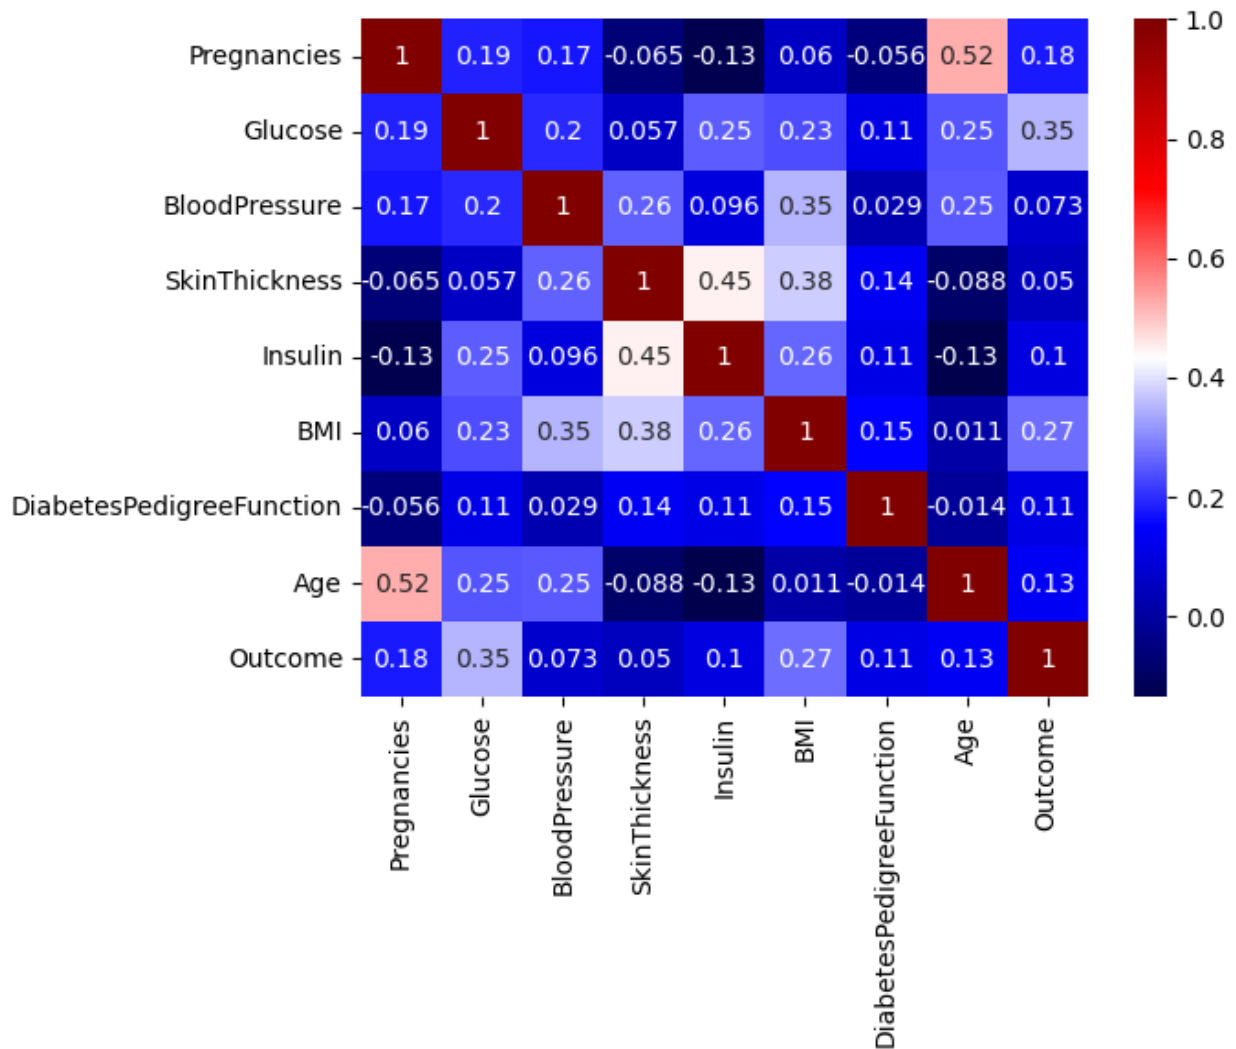
```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7a57c7c64430>
```



```
sns.heatmap(df.corr(), cmap="seismic",annot=True)
```

```
<Axes: >
```



The following features don't significantly impact our output parameters

- Blood Pressure
- Skin Thickness

```
def minmax_scaler(df, column):
    min=np.min(df[column])
    max=np.max(df[column])
    df[column]=(df[column]-min)/(max-min)

minmax_scaler(df, "Pregnancies")
minmax_scaler(df, "Glucose")
minmax_scaler(df, "BloodPressure")
minmax_scaler(df, "SkinThickness")
minmax_scaler(df, "Insulin")
minmax_scaler(df, "BMI")
minmax_scaler(df, "Age")
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	0.428571	0.636364	0.580357	0.546875	0.349451	0.605284
1	0.285714	0.272727	0.285714	0.671875	0.389011	0.448510
2	0.071429	0.769697	0.616071	0.453125	0.443956	0.668681
3	0.357143	0.684848	0.464286	0.343750	0.268132	0.406329
4	0.071429	0.618182	0.651786	0.703125	0.371429	0.723246

	DiabetesPedigreeFunction	Age	Outcome
0	0.780996	0.355556	1
1	0.335059	0.177778	1
2	0.618149	0.311111	1
3	0.147669	0.355556	0
4	0.078000	0.000000	0

```
#
X=df[['Pregnancies','Glucose','BloodPressure','SkinThickness','Insulin',
      'BMI','DiabetesPedigreeFunction','Age']]
X=df[['Pregnancies','Glucose','Insulin','BMI','DiabetesPedigreeFunction',
      'Age']]
y=df[['Outcome']]

X_train=X.iloc[:400]
y_train=y.iloc[:400]
X_test=X.iloc[400:]
y_test=y.iloc[400:]

X_train = X_train.to_numpy()
y_train = y_train.to_numpy().flatten() # Flatten the y_train to
ensure it's a 1D array
X_test = X_test.to_numpy()
y_test = y_test.to_numpy().flatten() # Flatten the y_test to
ensure it's a 1D array
```

3. Comparison of Stochastic Gradient Descent and Batch Gradient Descent using Linear Regression

Stochastic Gradient Descent

```
class SGD:
    def __init__(self):
        self.w = None
        self.b = None

    def predict(self, x):
        # Predicted value
        return np.dot(x, self.w) + self.b

    def mse(self, y, y_pred):
        mse = np.sum((y - y_pred) ** 2) / np.size(y)
        return mse

    def costf(self, y, f):
        n = len(y)
        J = np.sum((y - f) ** 2) / (2 * n)
        return J

    def modify(self, f, a, x, y):
        # a is the learning rate

        # Change 1: Select a random data point
        random_idx = random.randint(0, len(x) - 1)
        x_random = x[random_idx, :]
        y_random = y[random_idx]

        # Calculate the gradient for the selected data point
        gradient_w = x_random * (f[random_idx] - y_random)
        gradient_b = f[random_idx] - y_random

        # Update w and b using the stochastic gradient
        self.w = self.w - a * gradient_w
        self.b = self.b - a * gradient_b

    def model(self, x, y, learning_rate=0.1, max_iterations=1000,
threshold=0.01):
        n = x.shape[1]
        self.w = np.random.rand(n)
        self.b = random.random()
        f = self.predict(x)
```



```

J = self.costf(y, f)
c = 0
J_naught = threshold
mse_epochs=[]

while J > J_naught and max_iterations > 0:
    J_prev = J
    self.modify(f, learning_rate, x, y)
    f = self.predict(x)
    J = self.costf(y, f)
    mse_epochs.append(self.mse(y,f))
    if J_prev < J:
        c += 1
    if c >= 5:
        learning_rate *= 0.1
        self.w = np.random.rand(n)
        self.b = random.random()
        f = self.predict(x)
        J = self.costf(y, f)
        c = 0

    if J_prev > J and J_prev - J <= learning_rate * 10 ** (-
3):
        break
    max_iterations -= 1
return mse_epochs

def fit(self, x, y, learning_rate=0.1, max_iterations=1000,
threshold=0.01):
    return self.model(x, y, learning_rate, max_iterations,
threshold)

def linreg(self, x, y):
    y_predicted = self.predict(x)
    mse_value = self.mse(y, y_predicted)
    return mse_value, self.w, self.b

m=SGD()
mse_epochs_SGD=m.fit(X_train,y_train)

error,w,b=m.linreg(X_test,y_test)

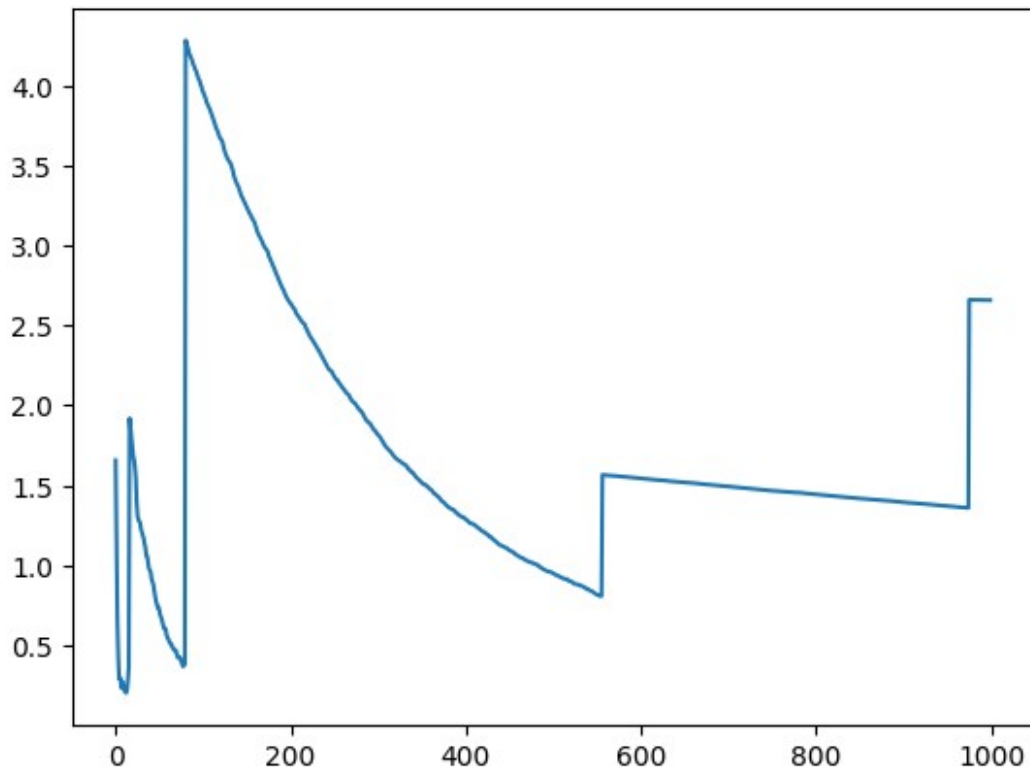
error

2.841988576015696

plt.plot(mse_epochs_SGD)

[<matplotlib.lines.Line2D at 0x7a57c304d000>]

```



Batch Gradient Descent

```
class BGD:
    def __init__(self):
        self.w = None
        self.b = None

    def predict(self, x):
        # Predicted value
        return np.dot(x, self.w) + self.b

    def mse(self, y, y_pred):
        mse = np.sum((y - y_pred) ** 2) / np.size(y)
        return mse

    def costf(self, y, f):
        n = len(y)
        J = np.sum((y - f) ** 2) / (2 * n)
        return J

    def modify(self, f, a, x, y):
        # Calculate the gradient using the entire dataset
        gradient_w = np.dot(x.T, (f - y)) / len(y)
        gradient_b = np.sum(f - y) / len(y)

        # Update w and b using the batch gradient
```

```

        self.w = self.w - a * gradient_w
        self.b = self.b - a * gradient_b

    def model(self, x, y, learning_rate=0.1, max_iterations=1000,
threshold=0.01):
        n = x.shape[1]
        self.w = np.random.rand(n)
        self.b = random.random()
        f = self.predict(x)
        J = self.costf(y, f)
        c = 0
        mse_epochs=[]
        J_naught = threshold

        while J > J_naught and max_iterations > 0:
            J_prev = J
            self.modify(f, learning_rate, x, y)
            f = self.predict(x)
            J = self.costf(y, f)
            mse_epochs.append(self.mse(y, f))
            if J_prev < J:
                c += 1
            if c >= 5:
                learning_rate *= 0.1
                self.w = np.random.rand(n)
                self.b = random.random()
                f = self.predict(x)
                J = self.costf(y, f)
                c = 0
            if J_prev > J and J_prev - J <= learning_rate * 10 ** (-
3):
                break
            max_iterations -= 1
        return mse_epochs

    def fit(self, x, y, learning_rate=0.1, max_iterations=100,
threshold=0.1):
        return self.model(x, y, learning_rate, max_iterations,
threshold)

    def linreg(self, x, y):
        y_predicted = self.predict(x)
        mse_value = self.mse(y, y_predicted)
        return mse_value, self.w, self.b

m=BGD()
mse_epochs_BGD=m.fit(X_train,y_train)

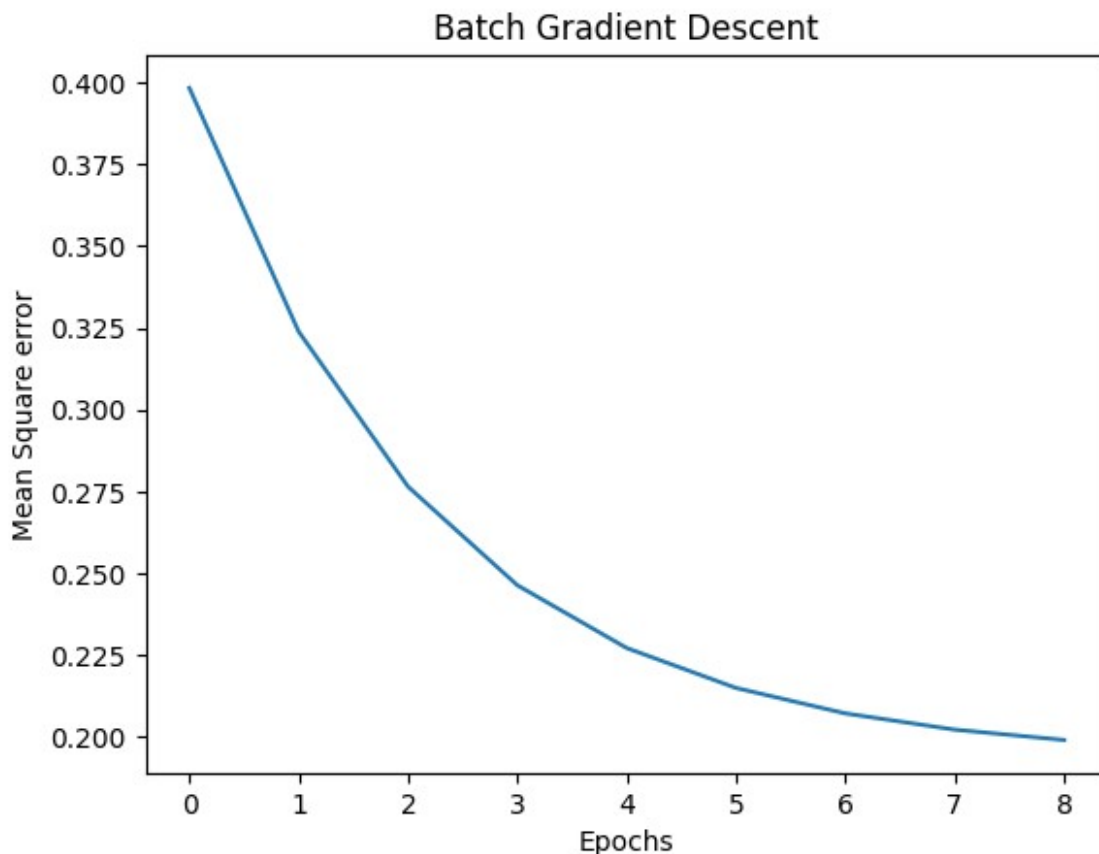
m.mse(y_test,m.predict(X_test))

```

0.2161486759430618

Insights drawn (plots, markdown explanations)

```
plt.plot(mse_epochs_BGD)
plt.xlabel('Epochs')
plt.ylabel('Mean Square error')
plt.title('Batch Gradient Descent')
Text(0.5, 1.0, 'Batch Gradient Descent')
```



Batch Gradient Descent Performs better than Stochastic Gradient Descent because the BGD operates on whole dataset to perform parameter tuning while the SGD takes a random data sample to do the same making it more prone to noise.

4. Comparison of Lasso and Ridge Regression using Polynomial Regression

Lasso Regression

```
class LassoRegression:
    def __init__(self, degree, lamda=0, learning_rate=0.01,
iterations=1000):
        self.degree = degree
        self.learning_rate = learning_rate
        self.iterations = iterations
        self.theta = None
        self.lamda=lamda

    def fit(self, X, y):
        X_poly = self.polynomial_features(X)
        self.theta = self.initialize(X_poly)
        mse_epochs=[]

        m = len(X)
        for _ in range(self.iterations):
            predictions = self.predict(X)
            error = predictions - y.reshape(-1,1)
            gradient = (1/m) * X_poly.T.dot(error)
            mse_epochs.append(self.mse(y,predictions))

            regularization_term = self.l1_regularization()
            gradient += regularization_term

            self.theta -= self.learning_rate * gradient
        return mse_epochs

    def polynomial_features(self, X):
        X_poly = np.c_[X]

        for i in range(2, self.degree + 1):
            X_poly = np.c_[X_poly, X**i]

        X_poly = np.c_[np.ones((X_poly.shape[0], 1)), X_poly]

        return X_poly

    def initialize(self,X):
        return np.random.rand(X.shape[1], 1)

    def predict(self, X):
        X_poly = self.polynomial_features(X)
        # print(X_poly.shape)
```

```

        # print(self.theta.shape)
        return X_poly.dot(self.theta)

    def l1_regularization(self):      #Lasso
        return self.lamda * np.sign(self.theta)

    def mse(self, y, y_pred):
        mse = np.sum((y - y_pred) ** 2) / np.size(y)
        return mse

m = LassoRegression(degree=2, lamda=0.1)
mse_epochs_lasso=m.fit(X_train, y_train)
m.mse(y_test,m.predict(X_test))

25.94669128442396

```

Ridge Regression

```

class RidgeRegression:
    def __init__(self, degree, lamda=0, learning_rate=0.01,
iterations=1000):
        self.degree = degree
        self.learning_rate = learning_rate
        self.iterations = iterations
        self.theta = None
        self.lamda=lamda

    def fit(self, X, y):
        X_poly = self.polynomial_features(X)
        self.theta = self.initialize(X_poly)
        mse_epochs=[]

        m = len(X)
        for _ in range(self.iterations):
            predictions = self.predict(X)
            error = predictions - y.reshape(-1,1)
            gradient = (1/m) * X_poly.T.dot(error)
            mse_epochs.append(self.mse(y,predictions))

            regularization_term = self.l2_regularization()
            gradient += regularization_term

            self.theta -= self.learning_rate * gradient
        return mse_epochs

    def polynomial_features(self, X):

```

```

X_poly = np.c_[X]

for i in range(2, self.degree + 1):
    X_poly = np.c_[X_poly, X**i]

X_poly = np.c_[np.ones((X_poly.shape[0], 1)), X_poly]

return X_poly

def initialize(self,X):
    return np.random.rand(X.shape[1], 1)

def predict(self, X):
    X_poly = self.polynomial_features(X)
    # print(X_poly.shape)
    # print(self.theta.shape)
    return X_poly.dot(self.theta)

def l2_regularization(self):    #Ridge
    return self.lamda * 2 *self.theta

def mse(self, y, y_pred):
    mse = np.sum((y - y_pred) ** 2) / np.size(y)
    return mse

m = RidgeRegression(degree=2, lamda=0.1)
mse_epochs_ridge=m.fit(X_train, y_train)
m.mse(y_test,m.predict(X_test))

24.936790456035787

```

Insights drawn (plots, markdown explanations)

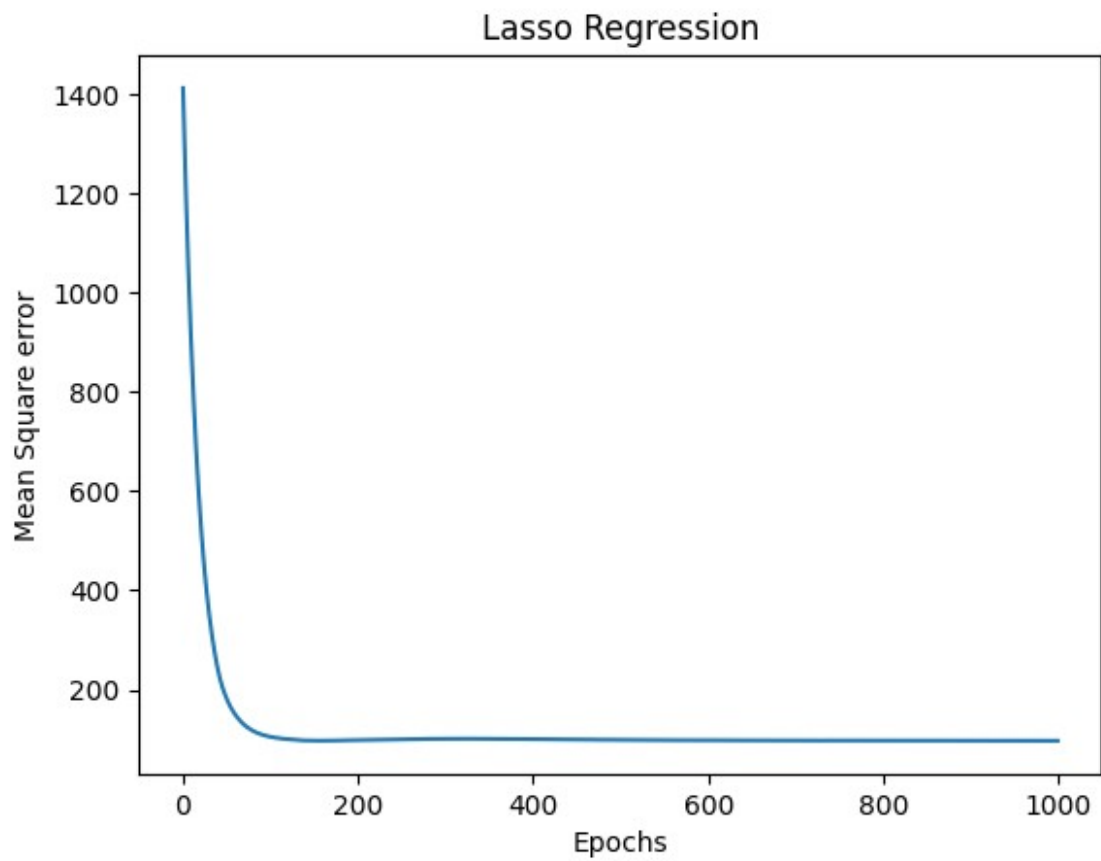
Lasso Regression

```

plt.plot(mse_epochs_lasso)
plt.xlabel('Epochs')
plt.ylabel('Mean Square error')
plt.title('Lasso Regression')

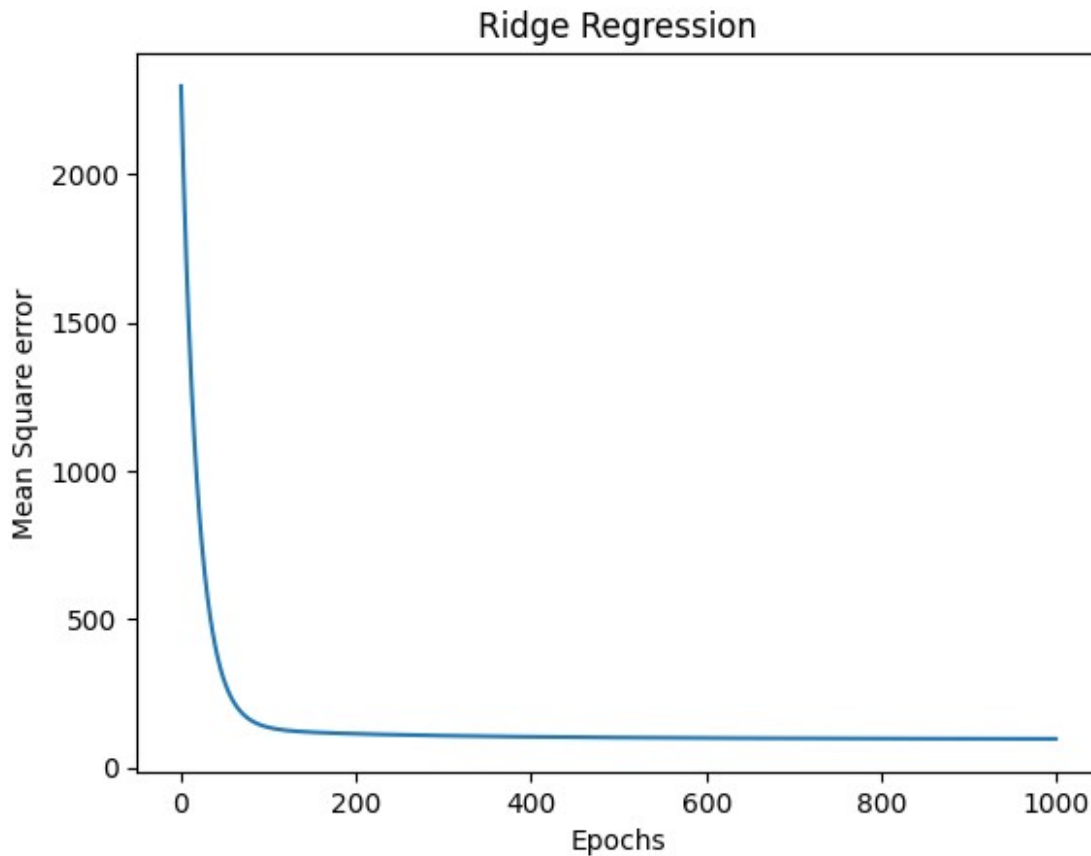
Text(0.5, 1.0, 'Lasso Regression')

```



Ridge Regression

```
plt.plot(mse_epochs_ridge)
plt.xlabel('Epochs')
plt.ylabel('Mean Square error')
plt.title('Ridge Regression')
Text(0.5, 1.0, 'Ridge Regression')
```

Lasso regularisation is better at reducing the mean square error as compared to ridge regularisation because in the former, all the redundant feature weights are eventually set to zero, while in the latter, the redundant weights still have an impact on the model as their weights never actually become exactly zero.

5. Comparison of Logistic Regression and Least Squares Classification

Logistic Regression

```
class LogisticRegression:

    def __init__(self, learning_rate=0.015, epochs=100000):
        self.learning_rate = learning_rate
        self.epochs = epochs
        self.weights = None
        self.bias = None

    def fit(self, X, y):
        samples, features = X.shape
```

```

# initialize parameters
self.weights = np.random.rand(features, 1)
self.bias = 0
loss=[]

# gradient descent
for _ in range(self.epochs):
    # approximate y with linear combination of weights and x,
plus bias
    linear_model = np.dot(X, self.weights) + self.bias
    # apply sigmoid function
    y_pred = self.sigmoid(linear_model)

    # compute gradients
    gradient_w = (1 / samples) * np.dot(X.T, (y_pred - y))
    gradient_b = (1 / samples) * np.sum(y_pred - y)
    # update parameters
    self.weights -= self.learning_rate * gradient_w
    self.bias -= self.learning_rate * gradient_b
    loss.append(self.loss(y,y_pred))

return loss

def predict(self, X):
    linear_model = np.dot(X, self.weights) + self.bias
    y_predicted = self.sigmoid(linear_model)
    y_pred_class = [1 if i > 0.5 else 0 for i in y_predicted]
    return np.array(y_pred_class)

def sigmoid(self, x):
    return 1 / (1 + np.exp(-x))

def loss(self,y,y_pred):
    return -np.mean(y * np.log(y_pred) + (1 - y) * np.log(1 -
y_pred))

def f1_score(self,y_true, y_pred):
    # Calculate precision and recall
    tp=0
    fp=0
    fn=0
    for i in range(len(y_true)):
        tp += (y_true[i] == 1) & (y_pred[i] == 1)
        fp += (y_true[i] == 0) & (y_pred[i] == 1)
        fn += (y_true[i] == 1) & (y_pred[i] == 0)

    precision = tp / (tp + fp) if (tp + fp) > 0 else 0
    recall = tp / (tp + fn) if (tp + fn) > 0 else 0

```

```

    # Calculate F1 score
    f1 = 2 * (precision * recall) / (precision + recall) if
(precision + recall) > 0 else 0

    return f1

def conf_mat(y_true, y_pred, threshold=0.5):
    y_pred_binary = (y_pred >= threshold).astype(int)

    TP = ((y_true == 1) & (y_pred_binary == 1)).sum()
    TN = ((y_true == 0) & (y_pred_binary == 0)).sum()
    FP = ((y_true == 0) & (y_pred_binary == 1)).sum()
    FN = ((y_true == 1) & (y_pred_binary == 0)).sum()

    return TP, TN, FP, FN

def cm_plot(y_test, y_pred):
    # Calculate TP, TN, FP, FN
    TP, TN, FP, FN = conf_mat(y_test, y_pred)
    confusion_matrix = np.array([[TN, FP], [FN, TP]])
    sns.heatmap(confusion_matrix, annot=True, fmt='d', cmap='Blues',
xticklabels=['Predicted 0', 'Predicted 1'], yticklabels=['Actual 0',
'Actual 1'])
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.title('Confusion Matrix')
    plt.show()

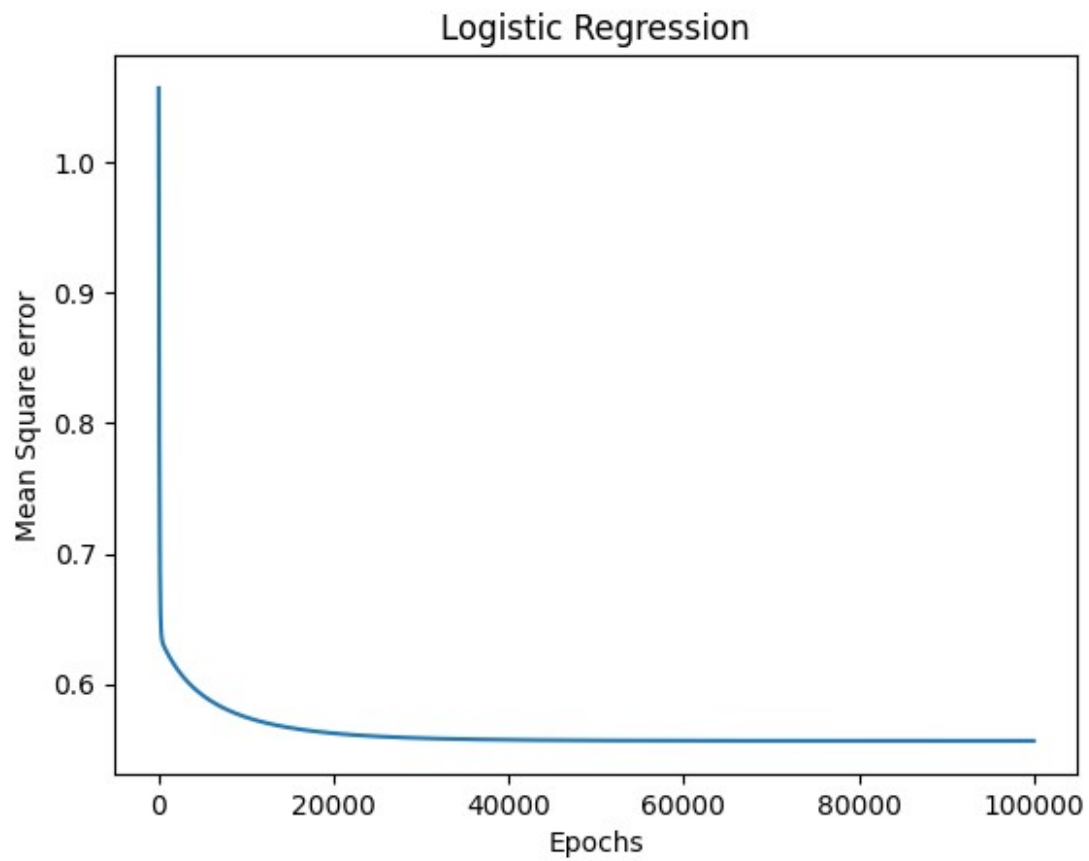
m=LogisticRegression()
y_train=y_train.reshape(-1,1)
y_test=y_test.reshape(-1,1)
loss=m.fit(X_train,y_train)
m.f1_score(y_test,m.predict(X_test))

array([0.50704225])

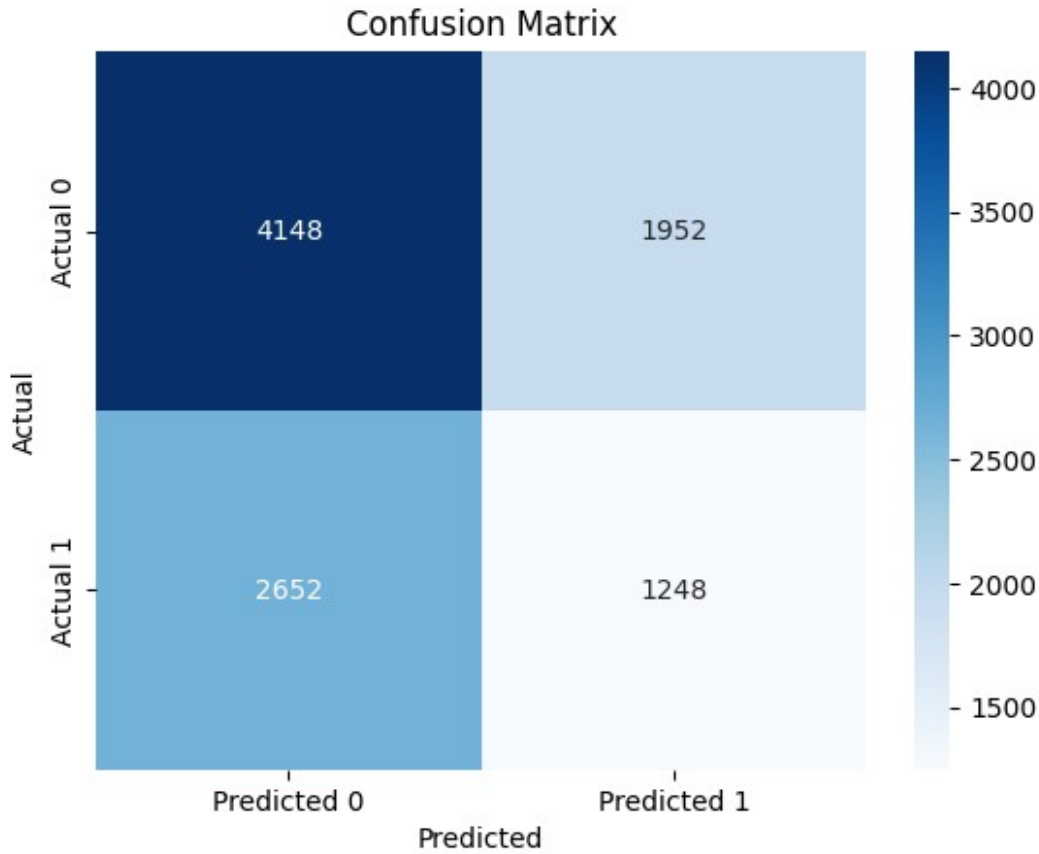
plt.plot(loss)
plt.xlabel('Epochs')
plt.ylabel('Mean Square error')
plt.title('Logistic Regression')

Text(0.5, 1.0, 'Logistic Regression')

```



```
y_pred=m.predict(X_test)  
cm_plot(y_test,y_pred)
```



Least Squares Classification

```
class LeastSquaresClassifier:
    def __init__(self):
        self.theta = None

    def fit(self, X, y):
        # Add a column of ones for the bias term
        X_b = np.c_[np.ones((X.shape[0], 1)), X]

        # Compute the least squares solution
        self.theta = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(y)

    def predict(self, X):
        # Add a column of ones for the bias term
        X_b = np.c_[np.ones((X.shape[0], 1)), X]

        # Make predictions
        y_predict = X_b.dot(self.theta)

        # Convert predictions to binary (0 or 1)
        return (y_predict > 0.5).astype(int)
```

```

def error(self,X,y):
    X_b = np.c_[np.ones((X.shape[0], 1)), X]
    return np.mean((y - X_b @ self.theta) ** 2)

def f1_score(self,y_true, y_pred):
    # Calculate precision and recall
    tp=0
    fp=0
    fn=0
    for i in range(len(y_true)):
        tp += (y_true[i] == 1) & (y_pred[i] == 1)
        fp += (y_true[i] == 0) & (y_pred[i] == 1)
        fn += (y_true[i] == 1) & (y_pred[i] == 0)

    precision = tp / (tp + fp) if (tp + fp) > 0 else 0
    recall = tp / (tp + fn) if (tp + fn) > 0 else 0

    # Calculate F1 score
    f1 = 2 * (precision * recall) / (precision + recall) if
    (precision + recall) > 0 else 0

    return f1

def plot_decision_boundary(self, X, y):
    plt.scatter(X[:, 1], y, color='blue') # Assuming X is a 2D
array and you want to use the second column
    plt.xlabel("X")
    plt.ylabel("y")

    # Plot the decision boundary
    x_decision = np.array([np.min(X[:, 1]), np.max(X[:, 1])]) #
Assuming X is a 2D array
    y_decision = -(self.theta[0] + self.theta[1] * x_decision) /
self.theta[2]
    plt.plot(x_decision, y_decision, "r-")

    plt.title("Least Squares Classification")
    plt.show()

m_LSC=LeastSquaresClassifier()
m_LSC.fit(X_train,y_train)
m_LSC.error(X_test,y_test)

0.19535087436507162

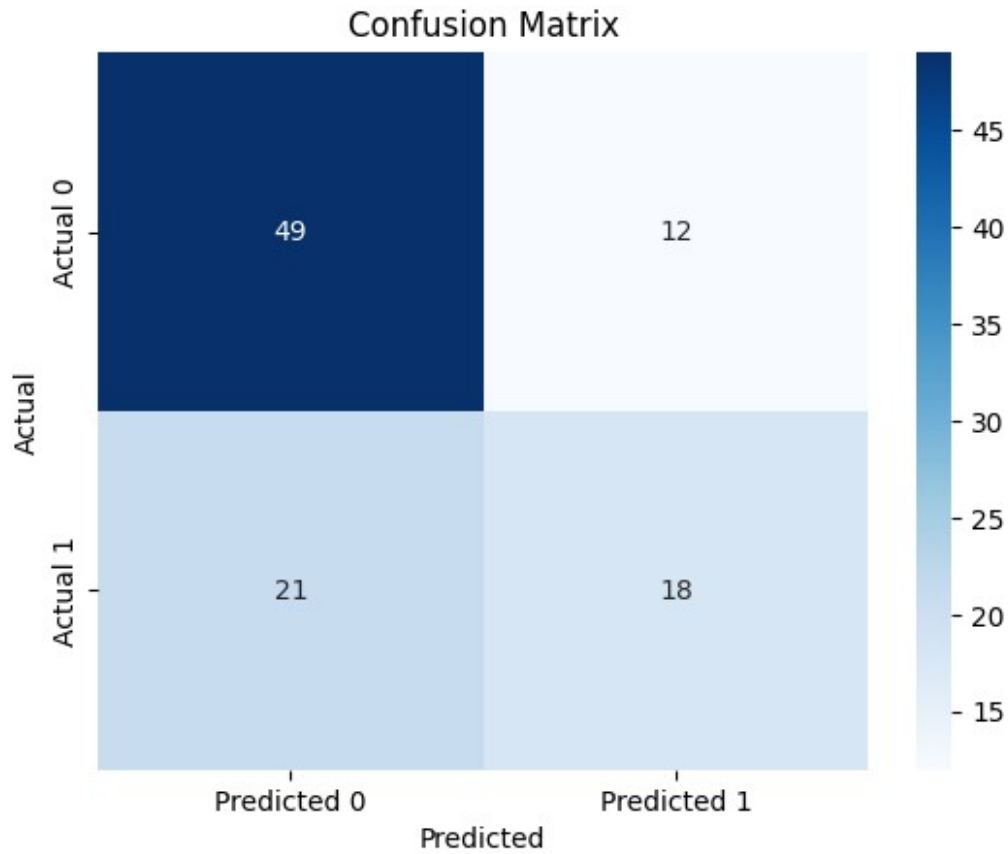
m_LSC.f1_score(y_test, m.predict(X_test))

array([0.50704225])

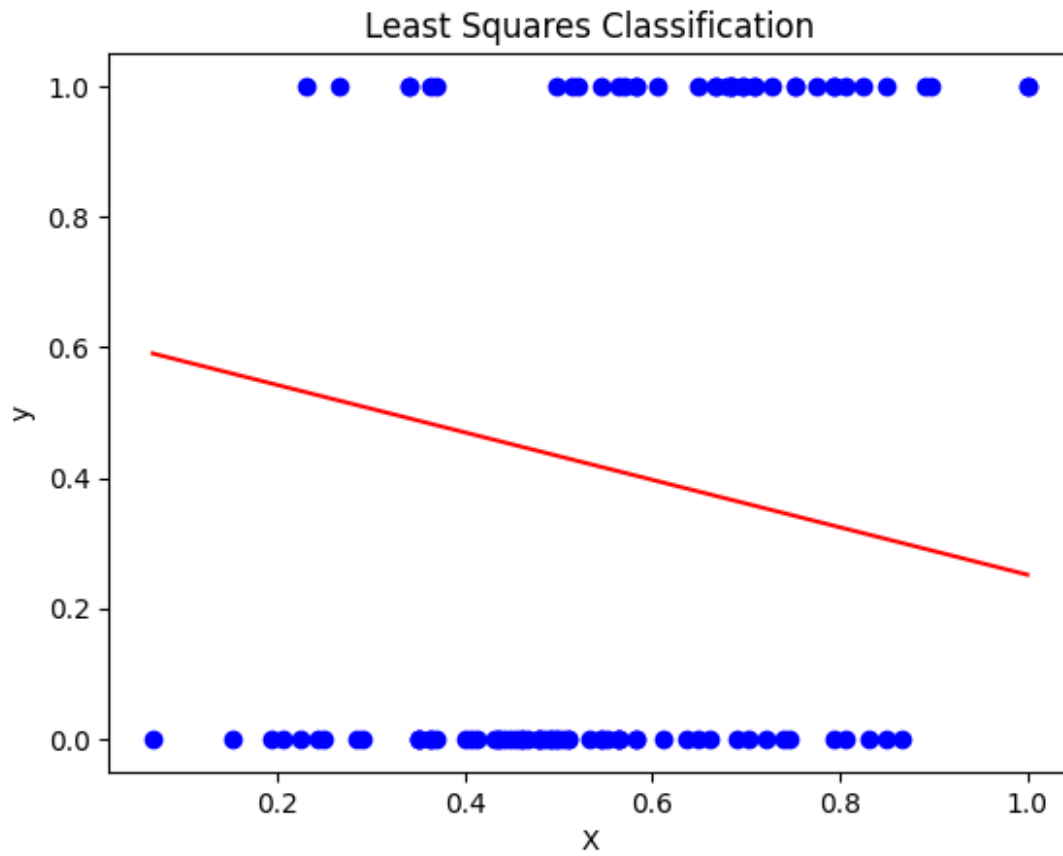
```

Insights drawn (plots, markdown explanations)

```
y_pred=m_LSC.predict(X_test)  
cm_plot(y_test,y_pred)
```



```
m_LSC.plot_decision_boundary( X_test, y_test)
```



Since the decision boundary required for classification in this scenario isn't too complex, a least squares classifier tends to perform better than a logistic classifier

6. References

1. <https://www.kaggle.com/code/jaypradipshah/polynomial-regression-from-scratch>
2. <https://www.kaggle.com/code/ninjaac/lasso-and-ridge-regression-from-scratch>
3. <https://marcoasilva.github.io/en/index.html>
4. <https://towardsdatascience.com/batch-mini-batch-stochastic-gradient-descent-7a62ecba642a>
5. <https://towardsdatascience.com/derivation-of-least-squares-regressor-and-classifier-708be1358fe9>
6. **Book - Christopher Bishop: Pattern Recognition and Machine Learning, Springer International Edition**